

Review of Related Literature and Studies

Prinz Edward L. Ducyogen

Centro Escolar University

Course Title

Professor Name

Due Date

Author Note

Correspondence concerning this article should be addressed to Prinz Edward L. Ducyogen, Email: ducyogenpred@gmail.com

Review of Related Literature and Studies

This chapter reviews related literature and studies that ground the proposed benchmarking of dynamic multi-task balancing on Taglish disaster tweets. Following the required thematic organization, the review arranges prior work by computational sub-concept, algorithmic paradigm, and empirical bottleneck rather than by geographical origin.

Section 2.1 examines crisis informatics corpora and the evidence for joint multi-task modeling of Named Entity Recognition (NER), intent classification, and urgency classification. Section 2.1 also examines the computational and statistical standards that such modeling must satisfy. Section 2.2 analyzes multi-task optimization (MTO) dynamics, covering dynamic loss-weighting methods, gradient-surgery and projection techniques, and the empirical evidence questioning whether specialized balancing algorithms outperform static scalarization. Section 2.3 reviews code-switched text processing, subword tokenization disparities, and Philippine language resources, including local disaster response studies. Section 2.4 consolidates key studies into an Algorithmic Literature Taxonomy Matrix positioned against the proposed research. Section 2.5 presents a three-step synthesis that derives the research gap and articulates how the proposed study bridges that gap.

The reviewed corpus draws on peer-reviewed sources indexed in the ACM Digital Library, IEEE Xplore, SpringerLink, ScienceDirect, and the Association for Computational Linguistics (ACL) Anthology. Sources were identified through database searches on multi-task optimization, crisis informatics, and code-switching conducted as

of 2026. Foundational multi-task optimization studies from 1997 to 2018 are retained where they define the algorithms under benchmark. The majority of the crisis informatics and code-switching sources that establish the evaluation domain date from the last five years.

2.1 Joint Multi-Task Modeling and Computational Efficiency in Crisis Informatics

2.1.1 Crisis Informatics Corpora and Benchmark Protocols

Crisis informatics research established that human-annotated social media corpora enable supervised processing of disaster messages at operational scales. Imran et al. (2016) collected over 52 million crisis-related tweets across 19 disaster events and released approximately 50,000 human-annotated messages. These messages spanned nine United Nations Office for the Coordination of Humanitarian Affairs (UN OCHA) humanitarian categories. Across Naive Bayes, Random Forest, and Support Vector Machine classifiers, the corpora achieved Area Under the ROC Curve (AUC) scores of at least 0.80 for most humanitarian categories. The rare “missing, trapped, or found people” category suffered measurably lower performance because of severe class imbalance. The authors additionally documented five categories of out-of-vocabulary noise in crisis microblogs, including misspellings, abbreviations, and phonetic substitutions, confirming that informal disaster text requires specialized lexical handling.

Subsequent benchmarks standardized label taxonomies, deduplication, and evaluation splits. Alam, Sajjad, et al. (2021) consolidated eight crisis tweet corpora into CrisisBench and demonstrated that duplicate and near-duplicate content is pervasive in uncleaned crisis datasets. Failing to remove such content therefore produces misleading

classification performance and risks data leakage between training and evaluation splits. On the consolidated humanitarian task, RoBERTa achieved a weighted F1-score of 0.872 compared with 0.829 for the best classical baseline. Preliminary multilingual experiments found that XLM-RoBERTa matched monolingual RoBERTa on the consolidated multilingual set, while systematic multilingual evaluation was left as future work. Alam, Qazi, et al. (2021) released HumAID, a 77,196-tweet corpus covering 19 disaster events between 2016 and 2019. The corpus was filtered from 24.8 million raw tweets through geographic, relevance, and near-duplicate pruning. Crowdsourced annotation reached an average Fleiss' kappa of 0.55. The HumAID benchmark showed that monolingual RoBERTa attained the highest average weighted F1-score of 0.781, followed by multilingual XLM-RoBERTa at 0.777, whereas classical classifiers remained below 0.72. Notably, DistilBERT matched full BERT performance (0.769 versus 0.768 average weighted F1) with substantially fewer parameters, a result the authors identified as directly relevant to latency-sensitive operational deployment.

Recent resources extend document-level corpora toward token-level annotations and multi-task formulations. Wang et al. (2021) formulated disaster triage as multi-task learning on the TREC Incident Streams corpus (40,459 tweets across 42 crisis events). The model jointly predicts 25 information types and a continuous priority score. The shared BERT encoder outperformed single-task learning on ranking quality (NDCG 0.5101 versus 0.4393) and alerting worth (-0.2689 versus -0.4057 for high-priority alerts). The encoder ablation showed that parameter-efficient variants (DistilBERT at 66 million parameters and ALBERT at 11 million parameters) remained competitive with full 110-million-parameter encoders. However, the study did not evaluate computational

efficiency or inference latency. Its fixed equal weighting is precisely the untested condition that Section 2.2 shows can suppress one task family in favor of another.

Token-level NER aggregates sequence-labeling gradients across every subword token, whereas sequence-level classification contributes a single pooled signal.

Gradient-magnitude disparity theory predicts that this asymmetry will degrade the weaker task’s gradients (Chen et al., 2018). No corpus in this lineage pairs token-level entity spans with concurrent sequence-level triage labels under cross-event evaluation, a void the proposed study fills directly.

2.1.2 Joint Versus Single-Task Modeling of Token-Level and Sequence-Level Tasks

Foundational work in multi-task learning (MTL) established the representational and statistical rationale for joint training. Caruana (1997) formalized hard parameter sharing, in which related tasks train in parallel over a common hidden representation. Caruana also demonstrated on road-following and medical risk prediction domains that the inductive transfer improves generalization on the primary task. Baxter (2000) proved statistically that learning a shared representation across K related tasks reduces the per-task sample complexity by a factor proportional to $1/K$. This reduction holds provided that tasks originate from a common environment distribution. These guarantees nevertheless presume task relatedness. When tasks conflict, joint training produces negative transfer that degrades the weaker task (Adouane & Bernardy, 2020). Predicting task compatibility and mitigating conflicting gradients therefore remain open problems. Sections 2.2 and 2.3 examine these two failure modes empirically.

Empirical studies of joint sequence labeling and sequence classification provide

direct precedent for the proposed NER, intent, and urgency task bundle. In task-oriented dialogue, Weld et al. (2022) surveyed joint intent detection and slot filling. Joint models consistently outperformed independent pipelines by capturing mutual conditional distributions and eliminating cascading error propagation. Slot F1 on the ATIS benchmark improved steadily across successive systems, and intent accuracy exceeded 99 percent. Weld et al. (2022) cautioned, however, that long-solved benchmarks increasingly measure dataset difficulty rather than language understanding. The survey identified standardized training budgets and out-of-domain generalization as open problems for the field. These findings jointly support the premise that pairing token-level NER with sequence-level intent and urgency heads offers strong mutual reinforcement. The benefits of pairing nevertheless depend on task relatedness, dataset scale, and optimization control.

2.1.3 Complexity Analysis, Statistical Rigor, and Reproducibility Infrastructure

The computational case for hard parameter sharing follows from asymptotic complexity analysis. Deploying K independent single-task models requires K forward passes per input and $O(K \cdot |\theta_{\text{backbone}}|)$ parameter storage. A shared encoder computes all task features in one forward pass, with task heads contributing less than one percent of total parameters. This architecture yields effective $\mathcal{O}(1)$ time and memory relative to task count K (Caruana, 1997). Lin and Zhang (2023) operationalized this architecture family in the LibMTL PyTorch library, which unifies 13 weighting strategies and 8 sharing architectures. Their NYUv2 benchmarks showed that exact gradient-balancing methods demand per-task gradients over all shared parameters and require representation-level approximations when parameter counts grow. The library initially lacked natural language

processing backbones.

Two methodological studies define the statistical standards that the proposed benchmark adopts. Reimers and Gurevych (2017) retrained published sequence-tagging systems across dozens of random seeds and showed that single reported scores are statistically misleading. Score variation across random seeds reached 1.01 F1 on CoNLL-2003 NER and 8.23 F1 on ACE 2005. Development-set selection correlated only weakly with test performance (Spearman $\rho = 0.229$). Demšar (2006) demonstrated that classifier comparisons across multiple datasets violate the normality and sphericity assumptions of parametric tests. The study recommended the Friedman test with the Iman-Davenport correction followed by post-hoc Nemenyi procedures as the default non-parametric protocol. These two studies jointly motivate the proposed design decisions of multi-seed evaluation under Leave-One-Event-Out (LOEO) cross-validation and two-tier non-parametric significance testing.

2.2 Multi-Task Optimization Dynamics: Balancing Conflicting Task Gradients

2.2.1 Gradient Surgery and Projection Techniques

Sener and Koltun (2018) reframed multi-task learning as multi-objective optimization and adapted the Multiple Gradient Descent Algorithm (MGDA) with a Frank-Wolfe line search to locate Pareto-stationary updates. Exact MGDA scales linearly with task count because the algorithm requires separate backward passes. The authors therefore proposed MGDA-UB, which solves the min-norm problem over representation gradients in a single backward pass. MGDA-UB accelerated training by a factor of 25 on the 40-task CelebA benchmark. Yu et al. (2020) identified the “tragic triad” of high

curvature, disparate gradient magnitudes, and conflicting directions, and proposed Projecting Conflicting Gradients (PCGrad). PCGrad projects each task gradient onto the normal plane of any conflicting gradient with negative cosine similarity. PCGrad raised Multi-task CIFAR-100 accuracy to 71.0 percent (versus 67.7 percent for independent training). It also solved all ten Meta-World MT10 manipulation tasks with multi-task Soft Actor-Critic where the vanilla joint agent failed half of the tasks. Ablations showed that altering both gradient directions and magnitudes is necessary. The authors acknowledged two structural costs. First, convergence can stall when cosine similarity approaches -1 . Second, computing per-task gradients multiplies backward passes linearly with task count.

B. Liu et al. (2021) strengthened the convergence guarantees with Conflict-Averse Gradient Descent (CAGrad). CAGrad maximizes the worst-case local task improvement within a trust region around the average gradient and provably converges to the average-loss optimum when the hyperparameter satisfies $c < 1$. CAGrad halved the average performance drop on CityScapes (11.64 percent versus 22.60 percent for shared MTAN under equal weighting) and matched or surpassed PCGrad across NYUv2 and Meta-World. A task-sampled variant delivered 2x to 5x speedups. L. Liu et al. (2021) approached fairness directly with Impartial MTL (IMTL). IMTL-G computes closed-form scaling factors so that the aggregated gradient makes equal projections onto every raw task gradient. IMTL-L rebalances loss scales through learnable parameters without distributional assumptions. Both variants remain invariant to manual loss rescaling and are computable through a single backward pass over shared representations. Navon et al. (2022) grounded gradient aggregation in cooperative bargaining theory, deriving Nash-MTL as the unique scale-invariant update satisfying Pareto optimality, symmetry,

and independence of irrelevant alternatives. Nash-MTL achieved the best overall trade-off on QM9 molecular regression (mean rank 2.5 versus 6.8 for linear scalarization, whose Δm reached 177.6 percent). On NYUv2, it reached a Δm of -4.04 percent versus $+0.20$ percent for CAGrad. However, each update requires solving a K -dimensional convex problem. The authors mitigated this cost by refreshing the bargaining solution only every 50 to 100 steps at a modest performance cost.

Recent methods continue to rederive balancing from structural analyses of gradients. Aligned-MTL minimizes the condition number of the gradient matrix through singular value decomposition (Senushkin et al., 2023). Similarity-Aware Momentum Gradient Surgery equalizes gradients when pairwise magnitude similarity signals conflict and modulates the first-order momentum otherwise (Borsani et al., 2025). However, neither variant has been evaluated on language tasks, let alone code-switched text. Both inherit the per-task backward-pass economics described above. Read together, the gradient-surgery family exhibits a common structural cost. PCGrad requires per-task backward passes, CAGrad adds a trust-region subproblem, and Nash-MTL solves a K -dimensional convex program per update. Each surgical intervention therefore adds training-time overhead that scales with task count. This cost is precisely what the empirical skepticism of Section 2.2.3 interrogates.

2.2.2 Dynamic Loss-Weighting Methods

An alternative algorithmic family adapts scalar loss weights without modifying gradient directions. Kendall et al. (2018) derived homoscedastic uncertainty weighting from a probabilistic likelihood in which each task receives a learnable variance parameter,

producing the loss term $\frac{1}{2\sigma_i^2}\mathcal{L}_i + \log \sigma_i$. On CityScapes, uniform weighting degraded semantic segmentation to 50.1 percent IoU, below the 59.4 percent single-task level. Learned uncertainty weighting recovered 63.4 percent IoU with a final weight ratio of 43:1:0.16 across segmentation, depth, and instance tasks. It outperformed grid-searched static weights (62.8 percent IoU) because the weights adapt during training. Chen et al. (2018) proposed GradNorm, which adjusts loss weights so that gradient norms at the final shared layer track relative inverse training rates. A single GradNorm run matched or surpassed the results of a 100-run weight grid search on NYUv2 (depth RMS error of 0.925 m at $\alpha = 1.5$). This run required approximately 5 percent computational overhead. S. Liu et al. (2019) introduced Dynamic Weight Average (DWA) within the Multi-Task Attention Network (MTAN). DWA balances tasks from prior-epoch loss descent rates alone and requires no additional backward passes. This design made training robust to the choice of weighting scheme because attention-based feature selection implicitly balances tasks. B. Liu et al. (2023) subsequently proposed FAMO, which equalizes task log-loss improvement rates through a first-order update on weighting logits. FAMO amortizes all gradient computations into a single backward pass, achieving $\mathcal{O}(1)$ time and space per iteration. FAMO outperformed linear scalarization on NYUv2 (Δm of -4.10 percent versus $+5.59$ percent) and CelebA (40 tasks). It trained 25 times faster than Nash-MTL on a two-task synthetic benchmark. The authors acknowledged reliance on temporal loss differences as directional-derivative estimates and the need to tune the regularization decay hyperparameter γ .

2.2.3 Empirical Skepticism: Scalarization Baselines and Optimizer Confounds

A converging body of evidence questions whether specialized MTO algorithms justify their costs when baselines are tuned properly. Kurin et al. (2022) benchmarked unitary (equal-weight) scalarization against MGDA, GradNorm, PCGrad, CAGrad, and IMTL across vision, reinforcement learning, and machine translation benchmarks under standardized hyperparameter budgets. Well-tuned equal weighting with Adam matched or exceeded every specialized method. The specialized algorithms ran up to 35 times slower because of their multiple backward passes. The authors interpreted many balancing schemes as implicit regularizers rather than superior optimizers. Xin et al. (2022) reached congruent conclusions on large-scale multilingual machine translation and vision. Dynamic task weights assigned by MTO algorithms remained nearly constant during training. Per-task gradient computation reduced training throughput from 11.52 to 4.81 steps per second. Reallocating that budget to model depth strictly improved the Pareto frontier. In addition, learning-rate underspecification produced six to seven times more score variance than seed variation, creating illusions of algorithmic superiority in undertuned comparisons.

Elich et al. (2024) identified the standard optimizer itself as the decisive confound. Across CityScapes, NYUv2, and CelebA, Adam outperformed SGD with momentum for every architecture and weighting scheme. The authors proved that Adam is partially loss-scale invariant for task-specific heads because loss scaling factors cancel between the first- and second-moment updates. Optimal uncertainty weighting is fully scale invariant. Under Adam, plain Equal Weighting (EW) achieved Pareto-optimal performance in three

of four benchmark configurations and was not dominated by any specialized MTO method trained with SGD. Most consequentially for diagnostic methodology, Elich et al. (2024) measured gradient cosine similarities between samples within a single task. Intra-task, inter-sample conflicts occurred as frequently as cross-task conflicts. Gradient magnitude disparities, not angular misalignment, constitute the primary optimization challenge under adaptive optimization. One recent audit nonetheless moderated this skepticism. Gama and Grassi Jr (2025) evaluated sixteen specialized optimizers against tuned uniform loss across progressively complex task configurations. Specialized optimizers surpassed uniform weighting once tasks grew sufficiently complex, although well-chosen fixed weights remained competitive. Collectively, these results indicate that claims of balancing gains must be established under optimizer-controlled, multi-seed, budget-matched conditions on the target domain. The proposed Taglish triage benchmark adopts these conditions by design. It extends the audit protocol of Elich et al. (2024) and Kurin et al. (2022) to code-switched crisis text.

2.3 Code-Switched Text Processing and Low-Resource Philippine NLP

2.3.1 Computational Challenges of Code-Switched Language

A systematic review documents the structural obstacles that confront code-switched natural language processing. Winata et al. (2023) reviewed more than 400 publications across three decades. The review documented severe resource concentration around English-paired bilingual data, dominance of language identification and sentiment tasks, and dataset attrition. Over 60 percent of surveyed corpora are no longer publicly accessible. Southeast Asian and Austronesian contact varieties, including Tagalog-English

code-switching (Taglish), remain critically underrepresented. Where Taglish evidence exists, models degrade along the mixing gradient. Adoptante et al. (2025) evaluated five open-source large language models in a retrieval-augmented answering pipeline for English, Tagalog, and Taglish queries. Every model performed best on English queries, next on Taglish, and worst on Tagalog. This ordering quantifies the additional difficulty that mixed-language inputs impose on generative models.

Subword tokenization quantifies a concrete computational penalty for underrepresented contact varieties. Petrov et al. (2023) measured token fertility across open multilingual tokenizers and showed systematically higher fertility for low-resource languages. The same study documented sequence-length inflation of up to 15x for non-Latin scripts, which prematurely exhausts context windows and increases inference latency. For a Taglish triage pipeline, these findings imply that subword over-fragmentation dilutes token-level embeddings for NER, where span labels must survive fragmentation across more subword pieces. Over-fragmentation also increases optimization variance precisely where supervision is densest. These findings motivate multilingual encoders whose vocabularies cover Tagalog-English mixing.

2.3.2 Multi-Task Learning on Code-Switched Text

Small but growing literature tests multi-task architectures on code-switched text, and the findings remain mixed. Most directly relevant to the proposed research, Adouane and Bernardy (2020) conducted controlled pairwise MTL experiments on low-resource, noisy, code-switched Algerian user-generated text. The tasks spanned code-switch detection, NER, spelling normalization, and sentiment analysis. The results demonstrated

that multi-task benefits are conditional rather than universal. Pairing NER with sentiment analysis induced negative transfer that collapsed NER macro F1 from 49.68 to 34.60. In contrast, pairing NER with code-switch detection produced mutual gains that raised detection macro F1 from 64.54 to 71.29. The authors further documented that joint-training outcomes were unstable across task pairings and dataset sizes. The mechanism most conspicuously absent from their experimental controls was an explicit loss-balancing procedure. This evidence establishes both the promise and the risk of the proposed three-task Taglish bundle. It supplies the strongest empirical motivation for benchmarking explicit MTO algorithms rather than assuming automatic joint-training benefits. It additionally motivates the proposed study's measurement of per-task-pair gradient cosine similarity. The Algerian results imply that the NER-intent and NER-urgency pairs may behave differently from the intent-urgency pair. Uncertainty weighting nevertheless remains the natural first balancing baseline for code-switched multi-task objectives ([Kendall et al., 2018](#)).

2.3.3 Philippine Language Resources and Local Disaster Response Studies

Philippine NLP has matured from vocabulary-scale models to benchmark infrastructure, but along a trajectory that bypasses disaster-domain multi-tasking. Cruz and Cheng ([2022](#)) pretrained RoBERTa Base (110M) and Large (330M) models on the TLUnified corpus. The new models improved test accuracy over prior Filipino BERT baselines by an average of 4.47 points (72.67 percent on election hate speech detection versus 68.60 percent). Scaling from Base to Large yielded only marginal gains (+0.09 to +0.90 points) because of pretraining corpus size. The evaluation was delimited to

classification without sequence labeling. Miranda (2023) addressed sequence labeling with TLUNIFIED-NER, a 208,247-token news corpus annotated for Person, Organization, and Location spans. The corpus has inter-annotator Cohen's kappa of 0.81 to 0.83 and reference F1 of 0.89. Fine-tuned supervised models consistently outperformed zero-shot prompting of commercial and open large language models. The corpus covers formal news text rather than informal code-switched social media. Montalan et al. (2025) introduced Batayan, an eight-task Filipino and Taglish benchmark of 3,800 natively validated test instances. The benchmark showed that Southeast Asian instruction-tuned models such as Gemma-SEA-LION-v3-9B reached an average macro F1 of 79.23 on natural language understanding tasks. Its evaluation is task-isolated and excludes the disaster domain. Herrera et al. (2022) contributed TweetTaglish, 21,150 code-switched tweets with token- and tweet-level language distribution annotations. The average composition comprises 55.2 percent Tagalog, 19.4 percent English, and 24.5 percent other tokens, validated at instance-level kappa of 0.70. Proof-of-concept regressors reached R^2 of 0.909 for English proportion prediction. The resource contains no triage annotations. On the encoder side, AI Singapore (2026) combined the ModernBERT architecture with a Gemma 3 SentencePiece tokenizer and pretrained on three trillion tokens spanning 13 Southeast Asian and regional languages, including Filipino. The combination improves tokenization fertility and compression rates for regional languages and supports longer context windows with greater computational efficiency. This resource anchors efficient encoder selection for the proposed Taglish triage backbone.

Local disaster-focused studies address single tasks on single events with non-Transformer models. Ermino et al. (2022) trained a lightweight feedforward artificial

neural network on 4,020 Typhoon Ulysses (Vamco) tweets for binary urgency evaluation. The model reported 99.17 percent validation accuracy and an F1 of 0.9882 on a 100-tweet confusion-matrix set. The authors acknowledged a fixed 8,000-word vocabulary and a single-event training distribution as limitations. At the cross-lingual frontier, Sarioglu Kayi et al. (2020) transferred urgency classification from English to Sinhala and Odia. The model reached 75.7 macro F1 on English with RoBERTa-large but only 63.5 and 62.6 zero-shot macro F1 on the target languages. The annotators conflated general status updates with actionable distress calls. Ray Chowdhury et al. (2020) aggregated more than 130,000 crisis tweets and applied Manifold Mixup to multilingual BERT. The approach improved out-of-event generalization across four held-out disaster test sets, including 79.36 macro F1 on the Philippines Flood test event. Per-class analysis showed Caution-and-Advice as a comparatively weak class at 70.3 F1. Cross-event generalization has since become an explicit evaluation target in its own right. Seeberger et al. (2025) showed that crisis classifiers exploit spurious event-specific cues and proposed a causal debiasing framework with domain-specific attention experts. The framework outperformed all baselines across three disaster classification tasks under temporal splits to unseen events. This line of evidence justifies the cross-event Leave-One-Event-Out design of the proposed study. Across this body of work, no study combines token-level entity extraction, intent categorization, and urgency prioritization on code-switched Philippine disaster text under cross-event evaluation.

2.4 Algorithmic Literature Taxonomy Matrix

Table 1 positions the principal related studies against the proposed research across algorithm, dataset or testbed, evaluated metrics, and identified limitations. The matrix organizes studies into four thematic clusters in sequence. Crisis corpora and multi-task triage come first. MTO algorithms follow. Code-switched and Philippine NLP studies come next. The proposed study closes the matrix. Each identified limitation states the concrete absence that the proposed study supplies.

Table 1

Algorithmic Literature Taxonomy Matrix

Author & Year	Algorithm/Model Used	Dataset/Testbed	Evaluated Metrics	Identified Limitation/Gap
Imran et al. (2016)	NB, RF, SVM over domain CBOW embeddings	19 disasters; 52M tweets, ~50K annotated	AUC, 10-fold cross-validation	English-only; class imbalance on rare categories; no token-level spans; no joint tasks
Alam, Sajjad, et al. (2021)	CNN, fastText, BERT, RoBERTa fine-tuning	CrisisBench consolidated crisis tweets	Weighted precision, recall, F1, accuracy	Single-label scheme; no entity or urgency annotations; multilingual evaluation preliminary
Alam, Qazi, et al. (2021)	fastText, BERT, DistilBERT, RoBERTa, XLM-R	HumAID: 77,196 tweets, 19 events	Weighted precision, recall, F1; inter-annotator agreement	English-only; mutually exclusive labels; kappa 0.55; no latency or memory reporting
Wang et al. (2021)	Shared BERT MTL (25-class info types + priority regression)	TREC-IS: 40,459 tweets, 42 events	NDCG, alert worth, F1, accuracy, harmonic mean	Equal weighting only; no gradient diagnostics; efficiency, latency, and throughput unmeasured

Author & Year	Algorithm/Model Used	Dataset/Testbed	Evaluated Metrics	Identified Limitation/Gap
Ray Chowdhury et al. (2020)	Multilingual BERT with class-masked loss and Manifold Mixup	130K+ multilingual crisis tweets from CrisisNLP and CrisisLex with held-out event test sets	Macro F1 per held-out event, zero-shot language transfer	Multi-label classification only; no token-level spans or urgency tiers; no balancing control; Caution-and-Advice class weak at 70.3 F1
Sarioglu Kayi et al. (2020)	Ensemble classifiers over RoBERTa, LASER, XLM-R, and projection embeddings with self-training	English urgency corpus (1,919 tweets); Sinhala and Odia evaluation sets	Macro F1 (75.7 English; 63.5 and 62.6 zero-shot)	Annotations conflated status updates with distress calls; no Tagalog or code-switching; single urgency task only
Ermino et al. (2022)	Feedforward artificial neural network	4,020 Typhoon Ulysses tweets (100-tweet confusion set)	Validation accuracy (99.17 percent), F1 (0.9882)	Fixed 8,000-word vocabulary; single-event distribution; binary urgency only; no NER, intent, or Transformer baselines
Seeberger et al. (2025)	Causal debiasing with event bias model, domain-specific attention experts, masking augmentation	HumAID, CrisisLex T26, TREC-IS under temporal event splits	Per-task performance against all baselines on three disaster tasks	Classification-only; no code-switched text; no token-level annotations; no latency or memory reporting
Sener and Koltun (2018)	MGDA-UB (Frank-Wolfe Pareto-stationary updates over representation gradients)	MultiMNIST, CelebA, CityScapes	Per-task accuracy versus uncertainty weighting, GradNorm, single-task	Pareto-stationary points arbitrary without task preferences; representation-gradient approximation; no NLP

Author & Year	Algorithm/Model Used	Dataset/Testbed	Metrics	Evaluated	Identified Limitation/Gap
Kendall et al. (2018)	Homoscedastic Uncertainty Weighting	CityScapes (vision)	Segmentation IoU, instance AP, centroid pixel error, depth RMSE, weight trajectories		Vision-only; homoscedastic noise assumption; no text or NLP evaluation
Chen et al. (2018)	GradNorm (gradient-norm targets via training rates)	NYUv2, synthetic regression (vision)	Task-normalized loss, depth RMSE, segmentation IoU, surface normal error, keypoint error		Balances only the final shared layer; ~5% overhead; no language tasks
S. Liu et al. (2019)	MTAN with Dynamic Weight Average (prior-epoch loss descent rates)	CityScapes, NYUv2, Visual Decathlon	mIoU, depth error, task accuracy across weighting schemes		Weights react to loss history only; vision-only; attention heads couple architecture to task set; no NLP
Yu et al. (2020)	PCGrad (orthogonal projection of conflicting gradients)	CIFAR-100, NYUv2, Meta-World MT10	Accuracy, classification error, segmentation mIoU, depth error, task success rate		K backward passes; order-dependent projections; stall risk at cosine -1 ; no NLP
B. Liu et al. (2021)	CAGrad (trust-region worst-case improvement)	NYUv2, CityScapes, MT10 (vision)	Training loss, per-task performance, mIoU, depth error, success rate, training time		Trust-region hyperparameter c requires per-dataset tuning; no NLP evaluation
L. Liu et al. (2021)	IMTL-G / IMTL-L (equal-angle projection + loss balance)	NYUv2, CityScapes, CelebA (vision)	mIoU, depth error, normals error, accuracy		Assumes equal task priority; vision-only; no gradient-conflict logging

Author & Year	Algorithm/Model Used	Dataset/Testbed	Evaluated Metrics	Identified Limitation/Gap
Navon et al. (2022)	Nash-MTL (Nash bargaining over task gradients)	NYUv2, CityScapes, QM9, MT10	Mean rank, Δm , per-task metrics, runtime	Per-step K -dimensional convex solver; cost mitigated by infrequent updates; no code-switched text
B. Liu et al. (2023)	FAMO (log-loss improvement equalization, $\mathcal{O}(1)$)	NYUv2, CityScapes, QM9, CelebA, MT10	Δm , mean rank, runtime vs scalarization	Decay hyperparameter γ requires tuning; temporal-difference approximation; no NLP
Senushkin et al. (2023)	Aligned-MTL (gradient-matrix condition-number minimization via SVD)	Synthetic two-task, CityScapes, NYUv2, Meta-World MT10	Per-task metrics versus prior MTO methods	Requires the full gradient matrix each step; vision and reinforcement learning only; no language tasks
Borsani et al. (2025)	SAM-GS (similarity-triggered gradient equalization and momentum modulation)	Two synthetic two-task problems, CelebA, NYUv2, CityScapes, MT10	Mean ranking and improvement score versus prior MTO methods	Per-pair gradient similarity computations; vision and reinforcement learning only; no NLP or code-switched text
Kurin et al. (2022)	Unitary scalarization vs MGDA, GradNorm, PCGrad, CAGrad, IMTL	NYUv2, CityScapes, MT10	Per-task metrics, wall-clock runtime, GPU memory	Specialized methods up to 35x slower; scope limited to moderately related tasks; no code-switched NLP
Xin et al. (2022)	Empirical audit of MTO vs static scalarization	WMT translation, CityScapes, CelebA	BLEU, mIoU, error rates, steps/sec (11.52 to 4.81)	Dynamic weights nearly static in practice; Pareto grid search prohibitive; no low-resource text

Author & Year	Algorithm/Model Used	Dataset/Testbed	Metrics	Identified Limitation/Gap
Elich et al. (2024)	Equal Weighting + Adam; loss-scale invariance proofs	CityScapes, NYUv2, CelebA	Δm , mIoU, depth errors, gradient cosine/magnitude statistics	Vision-only; gradient statistics subsampled; language tasks untested
Gama and Grassi Jr (2025)	Audit of 16 specialized optimizers versus tuned uniform loss	Multi-MNIST, CityScapes, QM9	Per-task metrics and rankings across complexity tiers	Complexity thresholds established on vision and chemistry only; no language tasks; no efficiency profiling
Adouane and Bernardy (2020)	Shared-CNN MTL across CSD, NER, SPELL, SA	Algerian Arabic user-generated text	Accuracy, macro/micro F1	No principled loss balancing; task pairing decisive (NER+SA macro F1 fell 49.68 to 34.60)
Cruz and Cheng (2022)	RoBERTa Base and Large pretrained on TLUnified	TLUnified corpus; Filipino hatespeech, dengue, and NLI benchmarks	Test accuracy across three classification tasks	Classification-only; no sequence labeling; Base-to-Large gains marginal (+0.09 to +0.90); no disaster domain
Herrera et al. (2022)	Dictionary language ID + regression benchmark (MLP, SVR)	TweetTaglish: 21,150 Taglish tweets	R^2 , RMSE, inter-annotator kappa	Language-distribution labels only; word-level kappa 0.15 from cognate noise; no triage tasks
Miranda (2023)	Transformer NER baselines on gold-standard news corpus	TLUNIFIED-NER: 208,247 tokens, three entity types	Entity F1 (transformers above 0.90), Cohen's kappa 0.81 to 0.83	Formal news text only; no code-switched social media; no intent or urgency labels
Montalan et al. (2025)	Eight-task Filipino and Taglish benchmark evaluation	Batayan: 3,800 natively validated test instances	Macro F1 per task; 79.23 average for Gemma-SEA-LION-v3-9B on NLU tasks	Task-isolated evaluation; excludes disaster domain; no multi-task optimization comparisons

Author & Year	Algorithm/Model Used	Dataset/Testbed	Evaluated Metrics	Identified Limitation/Gap
Adoptante et al. (2025)	Five open-source LLMs in a LangChain and FAISS retrieval-augmented pipeline	Curated English, Tagalog, and Taglish documents with 15 QA pairs per document	Overlap metrics and LLM-judged groundedness and relevance	Question answering only; small evaluation set; no triage annotations; no multi-task formulation
Proposed study	Dynamic loss weighting (Uncertainty Weighting, GradNorm, FAMO) and gradient surgery (PCGrad, CAGrad, IMTL-G, Nash-MTL) versus STL, Uniform Equal Weighting, and tuned Static Linear Scalarization on a shared multilingual Transformer encoder	Multi-Task Corpus of Taglish Disaster Tweets (4 typhoon events; 4-fold LOEO; 3 seeds; RTX 4060 8GB workstation testbed)	Macro F1 (intent, urgency), span F1 (NER), Δm , t_{exec} (B=1), throughput (B=32), peak VRAM/RAM, training runtime, gradient norms and pairwise cosine similarities	Addresses the literature gap directly: first MTO benchmark on code-switched Taglish disaster triage with concurrent token- and sequence-level annotations, gradient-conflict diagnostics against an intra-task half-batch baseline, cross-event validation, and latency and memory profiling under AdamW optimization

Note. MTL = multi-task learning. MTO = multi-task optimization. NER = named entity recognition. STL = single-task learning. CSD = code-switch detection. SPELL = spelling normalization. SA = sentiment analysis. NLI = natural language inference. QA = question answering. NDCG = normalized discounted cumulative gain. Δm = relative multi-task performance change.

2.5 Synthesis of Related Literature and Studies

Three established trends emerge across the reviewed literature. First, crisis informatics has converged on Transformer encoders as the state-of-the-art backbone for

humanitarian message processing. Monolingual RoBERTa-class models lead English benchmarks (weighted F1 of 0.872 on CrisisBench and 0.781 on HumAID).

DistilBERT-class variants demonstrate comparable accuracy at lower parameter counts for deployment-oriented settings ([Alam, Sajjad, et al., 2021](#); [Alam, Qazi, et al., 2021](#)).

Second, joint multi-task modeling of token-level and sequence-level tasks is an accepted methodology that outperforms independent pipelines when tasks are compatible.

Evidence includes intent detection and slot filling ([Weld et al., 2022](#)) and disaster categorization with shared encoders ([Wang et al., 2021](#)). Hard parameter sharing provides effective $\mathcal{O}(1)$ inference cost relative to task count ([Caruana, 1997](#)). Third, the MTO literature has produced two mature algorithmic families. The first is dynamic loss weighting ([Chen et al., 2018](#); [Kendall et al., 2018](#); [B. Liu et al., 2023](#)). The second is gradient surgery ([Borsani et al., 2025](#); [B. Liu et al., 2021](#); [L. Liu et al., 2021](#); [Navon et al., 2022](#); [Senushkin et al., 2023](#); [Yu et al., 2020](#)). Both families have been evaluated with increasing methodological rigor under multi-seed, non-parametric protocols ([Demšar, 2006](#); [Reimers & Gurevych, 2017](#)). For code-switched and low-resource text, the field additionally agrees on prerequisites for usable systems ([Miranda, 2023](#); [Petrov et al., 2023](#); [Winata et al., 2023](#)). The prerequisites are multilingual encoders, careful annotation validation, and tokenization-aware design.

Four unresolved bottlenecks persist at the intersection of these trends. First, the central empirical dispute in MTO remains unresolved. Tuned static scalarization with adaptive optimizers matches or exceeds specialized balancing algorithms on vision and translation benchmarks ([Elich et al., 2024](#); [Kurin et al., 2022](#); [Xin et al., 2022](#)). However, every audit study operated on well-resourced, monolingual, or vision data. The question

of whether dynamic balancing earns its runtime overhead on low-resource, code-switched crisis text therefore remains unanswered. Second, the disaster MTL studies closest to the proposed problem stop short of optimization control. Wang et al. (2021) applied fixed weighting without gradient diagnostics or efficiency measurements. The study did not report inference latency, GPU memory, or cross-event generalization. The token-versus-sequence optimization behavior of disaster MTL therefore remains undiagnosed. Third, no existing corpus combines token-level entity spans, intent categories, and urgency tiers on code-switched Taglish disaster text. Philippine resources are task-isolated, formal-domain, or language-ID-oriented (Cruz & Cheng, 2022; Herrera et al., 2022; Miranda, 2023; Montalan et al., 2025). Local disaster studies address single tasks on single events with non-Transformer models (Ermino et al., 2022). In addition, prior code-switched MTL evidence warns that unconditional joint training can collapse minority task performance (Adouane & Bernardy, 2020). Fourth, diagnostic confounding remains unaddressed in applied settings. Elich et al. (2024) showed that intra-task, inter-sample gradient conflicts rival cross-task conflicts. However, the decoupling procedure exists only in computer vision. No crisis informatics study separates genuine cross-task interference from mini-batch variance when justifying balancing overhead.

The proposed study bridges these gaps through four concrete contributions mapped directly onto the bottlenecks. To resolve the unresolved MTO dispute on the target domain, the study benchmarks seven candidate algorithms against isolated Single-Task Learning, Uniform Equal Weighting, and tuned Static Linear Scalarization. The seven candidates span both algorithmic families, dynamic loss weighting (Uncertainty Weighting, GradNorm, FAMO) and gradient surgery (PCGrad, CAGrad, IMTL-G,

Nash-MTL). The evaluation applies an equal hyperparameter search budget, AdamW optimization (Loshchilov & Hutter, 2019), and a 4-fold Leave-One-Event-Out protocol across three seeds. This design directly extends the audit conditions of Elich et al. (2024) and Kurin et al. (2022) to code-switched crisis text. To supply the missing optimization evidence in disaster MTL, the framework integrates routines for logging gradients in PyTorch. These routines record per-task gradient norms, norm ratios, and pairwise cosine similarities on shared encoder layers. The framework pairs them with an intra-task baseline derived from half-batch gradients. This baseline operationalizes the decoupling principle of Elich et al. (2024) for crisis triage and enables Spearman correlation analysis between transfer gains Δm and measured conflict dynamics. To close the dataset void, the study constructs the Multi-Task Corpus of Taglish Disaster Tweets. The corpus provides concurrent, expert-adjudicated annotations for token-level entity spans, intent categories, and urgency tiers across four historical Philippine typhoons. This corpus is the first to pair the sequence-labeling depth of TLUNIFIED-NER (Miranda, 2023) with the code-switched linguistic coverage of TweetTaglish (Herrera et al., 2022). Finally, to ground deployment relevance, the benchmark records single-instance latency, batched throughput, and peak GPU and host memory under FP16 execution on an 8GB workstation. These measurements convert the deferred efficiency questions of Wang et al. (2021) into measured quantities. The study also integrates the resulting model into a prototype web application for disaster triage as the empirical demonstration vehicle required by Design Science Research evaluation. Through these four bridges, the study converts an unresolved optimization question, an undiagnosed gradient asymmetry, a missing multilingual corpus, and an unmeasured deployment cost into a single

reproducible empirical benchmark. This benchmark targets multi-task crisis triage on Taglish disaster tweets.

References

- Adoptante, A. J. M., Castro, J. A. D. V., Medrana, M. L. B., Ocampo, A. P. B., Peramo, E. C., & Miranda, M. R. M. (2025). Benchmarking open-source large language models on code-switched Tagalog-English retrieval augmented generation. *Journal of Advances in Information Technology*.
<https://doi.org/10.12720/jait.16.2.233-242>
- Adouane, W., & Bernardy, J.-P. (2020). When is multi-task learning beneficial for low-resource noisy code-switched user-generated Algerian texts? *Proceedings of the 4th Workshop on Computational Approaches to Code Switching*.
<https://aclanthology.org/2020.calcs-1.3>
- AI Singapore, A. P. P. (2026). *SEA-LION-ModernBERT and embedding*. <https://huggingface.co/collections/aisingapore/sea-lion-modernbert-and-embedding>
- Alam, F., Qazi, U., Imran, M., & Ofli, F. (2021). HumAID: Human-annotated disaster incidents data from Twitter with deep learning benchmarks. *Proceedings of the International AAAI Conference on Web and Social Media (ICWSM 2021)*.
<https://doi.org/10.1609/icwsml.v15i1.18116>
- Alam, F., Sajjad, H., Imran, M., & Ofli, F. (2021). CrisisBench: Benchmarking crisis-related social media datasets for humanitarian information processing. *Proceedings of the International AAAI Conference on Web and Social Media (ICWSM 2021)*. <https://doi.org/10.1609/icwsml.v15i1.18115>

- Baxter, J. (2000). A model of inductive bias learning. *Journal of Artificial Intelligence Research*. <https://doi.org/10.1613/jair.731>
- Borsani, T., Rosani, A., Nicosia, G., & Di Fatta, G. (2025). Gradient similarity surgery in multi-task deep learning. *Proceedings of the European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML PKDD 2025)*. https://doi.org/10.1007/978-3-662-72243-5/_6
- Caruana, R. (1997). Multitask learning. *Machine Learning*. <https://doi.org/10.1023/A:1007379606734>
- Chen, Z., Badrinarayanan, V., Lee, C.-Y., & Rabinovich, A. (2018). GradNorm: Gradient normalization for adaptive loss balancing in deep multitask networks. *Proceedings of the 35th International Conference on Machine Learning (ICML 2018)*. <https://arxiv.org/abs/1711.02257>
- Cruz, J. C. B., & Cheng, C. (2022). Improving large-scale language models and resources for Filipino. *Proceedings of the Thirteenth Language Resources and Evaluation Conference*. <https://arxiv.org/abs/2111.06053>
- Demšar, J. (2006). Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*. <https://www.jmlr.org/papers/v7/demsar06a.html>
- Elich, C., Kirchdorfer, L., Köhler, J. M., & Schott, L. (2024). Examining common paradigms in multi-task learning. *46th DAGM German Conference on Pattern Recognition (GCPR 2024)*. https://doi.org/10.1007/978-3-031-85181-0/_9
- Ermino, R. P., Abacan, R. C. A., Diamante, N. G. C., Faca, K. M. N., & Juntereal, T. E. D. (2022). Evaluating urgency of Typhoon-Related tweets through sentiment analysis using artificial neural networks. *International Journal of Computing Sciences*

Research. <https://doi.org/10.25147/ijcsr.2017.001.1.80>

Gama, G. S., & Grassi Jr, V. (2025). *Uniform loss vs. Specialized optimization: A comparative analysis in multi-task learning*. <https://arxiv.org/abs/2505.10347>

Herrera, M., Aich, A., & Parde, N. (2022). TweetTaglish: A dataset for investigating Tagalog-English code-switching. *Proceedings of the Thirteenth Language Resources and Evaluation Conference*. <https://aclanthology.org/2022.lrec-1.225>

Imran, M., Mitra, P., & Castillo, C. (2016). Twitter as a lifeline: Human-annotated Twitter corpora for NLP of crisis-related messages. *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC 2016)*.
<https://doi.org/10.63317/43bfut6ipc3j>

Kendall, A., Gal, Y., & Cipolla, R. (2018). Multi-task learning using uncertainty to weigh losses for scene geometry and semantics. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR 2018)*.
<https://doi.org/10.1109/CVPR.2018.00781>

Kurin, V., De Palma, A., Kostrikov, I., Whiteson, S., & Kumar, M. P. (2022). In defense of the unitary scalarization for deep multi-task learning. *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*.
<https://arxiv.org/abs/2201.04122>

Lin, B., & Zhang, Y. (2023). LibMTL: A python library for deep multi-task learning. *Journal of Machine Learning Research*.
<https://www.jmlr.org/papers/v24/22-0347.html>

Liu, B., Feng, Y., Stone, P., & Liu, Q. (2023). FAMO: Fast adaptive multitask optimization. *Advances in Neural Information Processing Systems 36 (NeurIPS*

2023). <https://arxiv.org/abs/2306.03792>

Liu, B., Liu, X., Jin, X., Stone, P., & Liu, Q. (2021). Conflict-averse gradient descent for multi-task learning. *Advances in Neural Information Processing Systems 34*

(*NeurIPS 2021*). <https://arxiv.org/abs/2110.14048>

Liu, L., Li, Y., Kuang, Z., Xue, J.-H., Chen, Y., Yang, W., Liao, Q., & Zhang, W. (2021).

TOWARDS IMPARTIAL MULTI-TASK LEARNING. *International Conference on Learning Representations*. <https://openreview.net/forum?id=IMPnRXEWpvr>

Liu, S., Johns, E., & Davison, A. J. (2019). End-to-end multi-task learning with attention.

Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2019). <https://doi.org/10.1109/CVPR.2019.00197>

Loshchilov, I., & Hutter, F. (2019). Decoupled weight decay regularization. *International Conference on Learning Representations*.

<https://openreview.net/forum?id=Bkg6RiCqY7>

Miranda, L. J. V. (2023). Developing a named entity recognition dataset for Tagalog.

Proceedings of the First Workshop in South East Asian Language Processing. <https://doi.org/10.18653/v1/2023.sealp-1.2>

Montalan, J. R., Layacan, J. P., Africa, D. D., Flores, R. I., Lopez II, M. T., Magsajo, T. D.,

Cayabyab, A., & Tjhi, W. C. (2025). Batayan: A Filipino NLP benchmark for evaluating large language models. *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.

<https://doi.org/10.18653/v1/2025.acl-long.1509>

Navon, A., Shamsian, A., Achituve, I., Maron, H., Kawaguchi, K., Chechik, G., & Fetaya, E. (2022). Multi-task learning as a bargaining game. *International Conference on*

Machine Learning (ICML 2022). <https://arxiv.org/abs/2202.01017>

Petrov, A., La Malfa, E., Torr, P. H. S., & Bibi, A. (2023). Language model tokenizers introduce unfairness between languages. *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. <https://arxiv.org/abs/2305.15425>

Ray Chowdhury, J., Caragea, C., & Caragea, D. (2020). Cross-lingual disaster-related multi-label tweet classification with Manifold Mixup. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. <https://aclanthology.org/2020.acl-srw.39>

Reimers, N., & Gurevych, I. (2017). Reporting score distributions makes a difference: Performance study of LSTM-networks for sequence tagging. *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*. <https://doi.org/10.18653/v1/D17-1035>

Sarioglu Kayi, E., Nan, L., Qu, B., Diab, M., & McKeown, K. (2020). Detecting urgency status of crisis tweets: A transfer learning approach for low resource languages. *Proceedings of the 28th International Conference on Computational Linguistics*. <https://aclanthology.org/2020.coling-main.414>

Seeberger, P., Freisinger, S., Bocklet, T., & Riedhammer, K. (2025). Generalizing to unseen disaster events: A causal view. *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics*. <https://doi.org/10.18653/v1/2025.findings-ijcnlp.2>

Sener, O., & Koltun, V. (2018). Multi-task learning as multi-objective optimization. *Advances in Neural Information Processing Systems 31 (NeurIPS 2018)*.

<https://arxiv.org/abs/1810.04650>

Senushkin, D., Patakin, N., Kuznetsov, A., & Konushin, A. (2023). Independent component alignment for multi-task learning. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR 2023)*.

<https://doi.org/10.1109/CVPR52729.2023.01923>

Wang, C., Nulty, P., & Lillis, D. (2021). Transformer-based multi-task learning for disaster tweet categorisation. *Proceedings of the 18th International Conference on Information Systems for Crisis Response and Management (ISCRAM 2021)*.

<https://arxiv.org/abs/2110.08010>

Weld, H., Huang, X., Long, S., Poon, J., & Han, S. C. (2022). A survey of joint intent detection and slot-filling models in natural language understanding. *ACM Computing Surveys*. <https://doi.org/10.1145/3547138>

Winata, G. I., Aji, A. F., Yong, Z.-X., & Solorio, T. (2023). The decades progress on code-switching research in NLP: A systematic survey on trends and challenges. *Findings of the Association for Computational Linguistics: ACL 2023*.

<https://doi.org/10.18653/v1/2023.findings-acl.185>

Xin, D., Ghorbani, B., Garg, A., Firat, O., & Gilmer, J. (2022). Do current multi-task optimization methods in deep learning even help? *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*. <https://arxiv.org/abs/2209.11379>

Yu, T., Kumar, S., Gupta, A., Levine, S., Hausman, K., & Finn, C. (2020). Gradient surgery for multi-task learning. *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*. <https://arxiv.org/abs/2001.06782>